

Attachment 1 — Artificial Intelligence (AI) Use Compliance and Risk Management Plan

Frasier Digital, LLC — RFQ 75D301-26-Q-00146, Modernized VAERS Reporting Application

Completed on the Government-provided template (Attachment 1, posted 2026-08-18). Frasier Digital discloses two distinct AI uses and completes one card for each: **Card A** — the in-application intelligent completion assistance capability (a Government-approved application feature delivered under the PWS), and **Card B** — contractor-side AI-assisted software development tooling. The CDCH.10 eleven-element narrative plan (ai-plan.md) is retained as a companion appendix.

CARD A — In-Application Intelligent Completion Assistance

1. Background

Name of AI tool/system	Intelligent completion assistance layer of the Modernized VAERS Reporting Application (application feature built by Frasier Digital on CDC's enterprise Azure OpenAI service, EDAV)
Contractor using the AI	Frasier Digital, LLC
Subcontractor or service provider, if any	None. The AI service is CDC's existing FedRAMP-authorized Azure OpenAI offering (EDAV); no third-party AI service is introduced
Contract task/function supported	PWS §1.6–1.8: intelligent completion assistance, plain-language field guidance, answering user questions, narrative review for completeness
Describe how AI will be used in contract performance	An optional, user-facing assistance layer: contextual guidance and plain-language explanation of VAERS fields; answers to user questions drawn from CDC-approved help content; optional narrative-to-form suggestions in which every AI-proposed answer is individually reviewed and confirmed by the user and visibly tagged until verified. Deterministic branching logic and field validation are rules-based, not AI. The AI layer never submits a report, alters submitted data, makes any determination, or blocks a user from proceeding. Design subject to Government approval at the wireframe/design review
Is AI use required to perform the work, or optional?	☒ Required (the PWS requires intelligent completion assistance; generative AI is our proposed implementation, subject to CO/COR design approval) ☐ Optional ☐ Unknown
AI classification	☒ Generative AI ☒ Natural Language Processing ☐ Agentic AI ☐ Classical/Predictive ML ☐ Computer Vision ☐ Reinforcement Learning
Is this use limited to contractor internal operations?	☐ Yes ☒ No
Will the AI touch Government information, systems, identities, networks, or operations?	☒ Yes (deployed entirely within the CDC-managed Azure environment; processes user-entered report assistance content inside the authorization boundary) ☐ No ☐ Unknown

Was the AI system procured to support this particular contract per requirements in the Statement of Work?	No new AI procurement. The application uses CDC's existing enterprise Azure OpenAI (EDAV) service consistent with the SOW; no new AI services, licenses, or subscriptions are introduced
---	--

2. Who owns and supports this use?

Contractor AI owner	Joseph Frasier, Program Manager / Lead Architect
Contractor technical support contact	Ryan Daley, Senior Engineer
Contractor privacy/security contact	Joseph Frasier, Security & Compliance Lead
CDC/HHS program office using or receiving outputs	CDC VAERS program office per RFQ 75D301-26-Q-00146 (Immunization Safety Office)
Expected users	☒ Public users ☒ Beneficiaries, patients, applicants, regulated entities, or other external parties (healthcare providers) ☒ CDC/HHS staff ☐ Contractor staff only

3. Tool background

Tool or product name	Azure OpenAI Service via CDC EDAV; Frasier-built application assistance layer
Developer/vendor	Microsoft (service); OpenAI (models); Frasier Digital (application layer, delivered with Government rights per FAR 52.227-17)
Open-source, proprietary, Government-furnished, or other?	☒ Government-furnished (EDAV Azure OpenAI service) ☒ Contractor-developed (application layer; Government rights per FAR 52.227-17) ☐ Open-source ☐ Proprietary/commercial ☐ Subcontractor-provided
If the system uses an LLM or foundation model, what is the primary model?	A GPT-family model from the CDC EDAV Azure OpenAI catalog; specific model/version selected with the COR at design review, pinned and documented in the delivered system documentation
Who provides the LLM/foundation model?	Microsoft Azure OpenAI, within the CDC authorization boundary
What model options are available?	Models available in CDC's EDAV Azure OpenAI catalog; configurable without code change
Does it require an API key?	☒ Yes — managed identity/keys held entirely within the CDC environment; no external keys, no calls leaving the boundary
Attach or link, if available	☒ Model card (Microsoft-published, incorporated by reference and attached to delivered system documentation) ☒ Vendor security documentation (Azure FedRAMP package, via CDC) ☒ Audit-log documentation (delivered with SSP artifacts) ☒ Terms of service / privacy policy (per HHS enterprise agreement) ☐ System card ☐ Data card ☐ Evaluation summary

4. Where is it hosted and authorized?

Hosting environment	☒ FedRAMP-authorized cloud ☒ CDC/HHS environment (CDC-managed Azure) ☐ Contractor environment ☐ Commercial SaaS ☐ Local/offline
Where is the system hosted?	United States Azure regions under CDC tenancy (EDAV)
Is the product directly FedRAMP authorized or available through a FedRAMP-authorized instance?	☒ Yes ☐ No ☐ Not applicable ☐ Unknown
If no, is it undergoing FedRAMP authorization?	Not applicable

5. What data will be used?

Will the AI process, store, transmit, summarize, analyze, or generate Government information?	☒ Yes (within the CDC boundary only) ☐ No ☐ Unknown
Types of data involved	☒ Public information (CDC-approved help/FAQ corpus) ☒ PII ☒ PHI (user-authored narrative content in a vaccine adverse-event context may contain health information; design minimizes — no identifiers are required to use the assistant, with PII-redaction filtering as defense-in-depth) ☐ Government nonpublic ☐ Procurement-sensitive ☐ Proprietary
What information will be entered into the AI tool?	User-typed questions; optional user-authored narrative drafts; non-identifying field context needed for guidance. All within the CDC boundary
What outputs will the AI generate?	Plain-language guidance and explanations; suggested field values (each individually user-reviewed, confirmed, and tagged until verified); missing-detail prompts for narrative completeness
Will prompts, uploads, outputs, logs, or user interactions be stored?	☒ Yes — logged within the CDC Azure environment only, fully accessible to CDC
Retention period	Per CDC records schedules and the approved SSP
Will Government data, prompts, outputs, or logs be used to train, fine-tune, improve, or benchmark any model?	☒ No
Will any data, prompts, outputs, or logs be shared with other users, tenants, vendors, subcontractors, or third parties?	☒ No
Explain storage, use, retention, deletion, and sharing	All Government data, prompts, outputs, and logs remain CDC property inside the CDC authorization boundary; accessible to CDC on demand; retained and deleted per CDC records schedules; never used for model training; never transferred outside the boundary. No limitations proposed on Government rights, access, or reuse

6. Is this high-impact?

Could the AI output materially affect an individual's or organization's rights, benefits, services, health, safety, access, eligibility, enforcement status, or other significant interests?	☒ No
Could CDC/HHS personnel rely on the AI output as a principal basis for a decision or action?	☒ No — human submitters author and submit all reports; CDC personnel perform all downstream review of reports, not AI outputs
Is the AI used only for internal drafting, summarization, or administrative support with human review before use?	☐ Yes ☒ No (public-facing assistance; every suggestion is human-reviewed by the submitter before any use, and the non-AI path is always complete)
High-impact rationale	Assessed not high-impact (OMB M-25-21 App. A): the AI provides optional informational assistance for voluntary form completion; it does not control or materially influence decisions affecting rights, benefits, safety, or access; assistance is equivalent across public and provider paths and fully accessible; no user is required to interact with AI to submit a report. Descriptive information supporting the Government's own assessment is provided throughout; we will support any required AI impact assessment

7. What safeguards are in place?

How will contractor personnel verify AI outputs before relying on them?	No AI output enters a report without the submitter's individual review and confirmation; applied suggestions carry a visible "AI-suggested — verify" marking until confirmed or edited. Retrieval-based answers draw only from the CDC-approved content corpus
What human review is required before AI-generated content is submitted to CDC/HHS?	The human submitter authors and submits every report. Assistant content configuration (help corpus, prompt templates, guardrails) passes CDC program-staff approval through the low-code content workflow before deployment
How will the contractor prevent unauthorized disclosure of Government information?	All processing inside the CDC FedRAMP boundary under ATD/ATO controls; role-based access; FIPS 140-validated encryption in transit and at rest; no external service calls
How will the contractor prevent misuse, unauthorized use, or corruption of the AI system?	Input sanitization and prompt-injection guardrails; a classification stage routes safety-critical inputs (medical emergencies, requests for medical advice, causality questions, shared personal information, injection attempts) to fixed approved responses before any generative model is involved; output filtering under a strict cannot-say policy; rate limiting; kill-switch configuration flag letting CDC disable the assist layer without redeployment

How will the contractor protect privacy, civil rights, and civil liberties?	Assistance is optional and equivalent across paths; the non-AI path is always complete; 508/WCAG conformance within the ACR scope; no identifiers required to use the assistant; PII-redaction filter on prompts; no scores, rankings, or determinations about any individual
What internal policies govern this AI use?	Frasier Digital AI-use policy: transparency, accountability, data integrity, accuracy, sensitive-information protection, foreseeable-risk review; named accountable owner; documented review gates; prohibition on Government data in unapproved AI tools

8. Monitoring and misuse protection

What is monitored?	☒ Accuracy/performance ☒ Performance degradation or drift ☒ Data leakage signals ☒ Abuse/misuse ☒ Prompt injection ☒ Unauthorized access
Does CDC/HHS have access to monitoring information, audit logs, or summary reports?	☒ Yes — logs live in the CDC environment and are fully CDC-accessible; sampled transcript reviews conducted jointly with CDC; satisfaction-survey signal monitored across submitter types
How are issues detected, escalated, and corrected?	Automated monitoring and periodic sampled review detect issues; findings enter corrective action under the QCP; escalation through the PM to the COR
When will CDC/HHS be informed?	Incidents reported through the same 1-hour CDC incident-reporting channel as all other events; anticipated AI changes reported to CO/COR before incorporation
What is the standard operating procedure for incidents or unacceptable outputs?	Immediate containment via the kill-switch flag (assist layer disabled without redeployment; form remains fully usable); incident report within 1 hour; root-cause analysis and corrective action under the QCP; Government approval before re-enable

9. Required contractor affirmation

Frasier Digital, LLC affirms that: AI will not be used unless approved by the Contracting Officer; the contractor has disclosed all planned or potential AI use; the contractor will not enter Government information into unapproved AI tools; the contractor will not use Government data to train, fine-tune, or improve an AI model unless expressly authorized; the contractor will maintain and update this AI Use Card throughout contract performance; the contractor will notify the Contracting Officer and COR before changed, expanded, or new AI use; the contractor will prevent unlawful discriminatory use or outputs; and the contractor will comply with applicable privacy, security, data rights, records, civil rights, and AI requirements.

CARD B — Contractor-Side AI-Assisted Software Development Tooling

1. Background

Name of AI tool/system	Commercial AI development assistants (Anthropic Claude / Claude Code and comparable code-assistance tools)
Contractor using the AI	Frasier Digital, LLC
Subcontractor or service provider, if any	None

Contract task/function supported	Software development under the PWS: code drafting, test generation, documentation drafting, 508 remediation suggestions
Describe how AI will be used in contract performance	Contractor engineers use AI assistants during development. All AI-assisted work product passes our standard review gates: author ≠ reviewer, senior-engineer review before merge, CI test and accessibility gates. Development uses synthetic data and public schemas only; CDC data never enters these tools
Is AI use required to perform the work, or optional?	☒ Optional (internal productivity tooling) ☐ Required ☐ Unknown
AI classification	☒ Generative AI ☒ Natural Language Processing ☐ Others
Is this use limited to contractor internal operations?	☒ Yes ☐ No
Will the AI touch Government information, systems, identities, networks, or operations?	☒ No (synthetic data and public schemas only; no Government information, credentials, or system access) ☐ Yes ☐ Unknown
Was the AI system procured to support this particular contract per requirements in the Statement of Work?	No — contractor-furnished internal tooling, not a contract deliverable or Government procurement

2. Who owns and supports this use?

Contractor AI owner	Joseph Frasier, Program Manager / Lead Architect
Contractor technical support contact	Ryan Daley, Senior Engineer
Contractor privacy/security contact	Joseph Frasier, Security & Compliance Lead
CDC/HHS program office using or receiving outputs	None directly; outputs are contractor work product delivered through normal review-gated deliverables
Expected users	☒ Contractor staff only

3. Tool background

Tool or product name	Anthropic Claude / Claude Code (primary); comparable commercial code-assistance tools
Developer/vendor	Anthropic (primary); other established commercial vendors as adopted under policy
Open-source, proprietary, Government-furnished, or other?	☒ Proprietary/commercial (contractor-licensed)
If the system uses an LLM or foundation model, what is the primary model?	Anthropic Claude model family (current production models)
Who provides the LLM/foundation model?	Anthropic
What model options are available?	Vendor production model catalog

Does it require an API key?	☒ Yes — contractor-held accounts/keys; no Government credentials involved
Attach or link, if available	☒ Terms of service ☒ Privacy/data retention policy (vendor-published, available on request)

4. Where is it hosted and authorized?

Hosting environment	☒ Commercial SaaS (contractor-side)
Where is the system hosted?	Vendor-operated United States cloud infrastructure
Is the product directly FedRAMP authorized or available through a FedRAMP-authorized instance?	☒ Not applicable — no Government information enters these tools
If no, is it undergoing FedRAMP authorization?	Not applicable

5. What data will be used?

Will the AI process, store, transmit, summarize, analyze, or generate Government information?	☒ No
Types of data involved	☒ Public and non-sensitive information (contractor-developed application code — no Government data; public schemas, synthetic test data only)
What information will be entered into the AI tool?	Application source code (contractor-developed, no Government data), synthetic data, public documentation
What outputs will the AI generate?	Draft code, tests, and documentation — all reviewed by the responsible senior engineer before commit
Will prompts, uploads, outputs, logs, or user interactions be stored?	☒ Yes — vendor-side per vendor retention policy; contains no Government information
Retention period	Per vendor policy; contractor accounts configured for no-training/limited-retention options where offered
Will Government data, prompts, outputs, or logs be used to train, fine-tune, improve, or benchmark any model?	☒ No Government data is present to be used
Will any data, prompts, outputs, or logs be shared with other users, tenants, vendors, subcontractors, or third parties?	☒ No
Explain storage, use, retention, deletion, and sharing	The controlling safeguard is data separation: CDC data never leaves the CDC environment and never enters contractor AI tooling. Tool-side storage therefore contains only open-source code and synthetic material

6. Is this high-impact?

Could the AI output materially affect an individual's or organization's rights, benefits, services, health, safety, access, eligibility, enforcement status, or other significant interests?	☒ No
Could CDC/HHS personnel rely on the AI output as a principal basis for a decision or action?	☒ No
Is the AI used only for internal drafting, summarization, or administrative support with human review before use?	☒ Yes
High-impact rationale	Not high-impact: internal development assistance with mandatory human review; no Government data, no decisions about individuals

7. What safeguards are in place?

How will contractor personnel verify AI outputs before relying on them?	Author ≠ reviewer; senior-engineer review before merge; CI test, lint, and accessibility gates; characterization tests for AI-drafted changes
What human review is required before AI-generated content is submitted to CDC/HHS?	All deliverables pass the same review gates regardless of drafting method; the responsible senior engineer owns every commit
How will the contractor prevent unauthorized disclosure of Government information?	Hard prohibition on Government data in contractor AI tools (Frasier AI-use policy); development on synthetic data and public schemas only
How will the contractor prevent misuse, unauthorized use, or corruption of the AI system?	Contractor-managed accounts, access control, and policy training; periodic policy compliance review
How will the contractor protect privacy, civil rights, and civil liberties?	No personal data enters the tools; deliverable-level protections are governed by Card A and the SSP
What internal policies govern this AI use?	Frasier Digital AI-use policy (as in Card A §7)


8. Monitoring and misuse protection

What is monitored?	☒ Abuse/misuse (policy compliance) ☒ Data leakage signals (repo scanning for secrets/sensitive content)
Does CDC/HHS have access to monitoring information, audit logs, or summary reports?	☒ Upon request
How are issues detected, escalated, and corrected?	Repo scanning and review gates; escalation to the PM; corrective action under the QCP
When will CDC/HHS be informed?	Any suspected touch of Government information by an unapproved tool is reported through the 1-hour incident channel

What is the standard operating procedure for incidents or unacceptable outputs?	Contain (revoke access/rotate credentials), report within 1 hour, root-cause, corrective action
---	---

9. Required contractor affirmation

Card A §9 affirmation applies in full to this use.

Signature: 

Joseph Frasier, Founder & Managing Member, Frasier Digital, LLC Date: 9-3-26